

Machine Learning Project

Objective

In this project, you will work on building, optimizing, and deploying a machine learning model using **Azure Machine Learning**, with optional integration of **Optuna** for hyperparameter tuning, **MLflow** for experiment tracking, and **Streamlit** or **Azure Web App** for creating a front-end interface for model inference.

This end-to-end machine learning pipeline will allow you to gain hands-on experience in model training, deployment, and creating user interfaces for real-world applications.

Project Description

You will:

1. **Choose a dataset:** Select a dataset from platforms like **Kaggle**, **UCI Machine Learning Repository**, or any other public dataset of your choice. The dataset should be appropriate for a classification, regression, or clustering problem.
2. **Preprocess and Explore the Data:** Perform necessary data preprocessing such as handling missing values, encoding categorical variables, and feature scaling.
3. **Build a Machine Learning Model:** Develop an initial machine learning model using algorithms like Random Forest, XGBoost, or others, depending on the nature of the problem.
4. **Optimize the Model (Optional):**
 - Use **Optuna** for hyperparameter optimization, which can search for the best combination of hyperparameters to maximize model performance.
 - Alternatively, you can use **Azure ML's HyperDrive** if you prefer a built-in solution for hyperparameter tuning.
5. **Track Experiments with MLflow:** Use **MLflow** for tracking the training process, storing metrics, and logging model parameters.
6. **Deploy the Model:**
 - **Deploy as a REST API:** Deploy your trained model as a real-time web service using **Azure Machine Learning**.
 - **Optional Front-End:** Build an interactive front-end application to make predictions using the deployed model:
 - **Streamlit:** Create a simple Streamlit app that takes user input through a form, sends it to the Azure ML endpoint, and displays the predicted results.
 - **Azure Web App:** Alternatively, you can use an **Azure Web App** to create a web interface for model inference, allowing users to input data and receive predictions.

7. **Monitor and Evaluate:** After deployment, ensure that you monitor the model's performance, track metrics like latency and accuracy, and log results using MLflow.
-

Project Tasks

1. **Step 1: Setup Azure ML Workspace**
 - Create an **Azure Machine Learning workspace** in the Azure Portal.
 - Set up a **Compute Instance** for development and a **Compute Cluster** for model training (if necessary).
 2. **Step 2: Data Selection and Preparation**
 - Choose a dataset from **Kaggle**, **UCI**, or another public repository.
 - Preprocess the data: handle missing values, encode categorical features, scale numerical features, and perform any necessary transformations.
 3. **Step 3: Model Training and Optimization**
 - Train an initial machine learning model using the dataset.
 - Optionally, use **Optuna** to tune hyperparameters for better model performance.
 - Log your training process using **MLflow**, capturing parameters, metrics, and models.
 4. **Step 4: Model Evaluation**
 - Evaluate your model using appropriate metrics such as accuracy, precision, recall, AUC, or RMSE, depending on the type of problem.
 - Fine-tune the model based on evaluation results.
 5. **Step 5: Model Deployment**
 - Register your best-performing model in the **Azure ML model registry**.
 - Deploy the model as an Azure ML **real-time endpoint**.
 6. **Step 6: Front-End Interface (Optional)**
 - Build a **Streamlit** or **Azure Web App** interface to allow users to input data and get predictions from the deployed model:
 - **Streamlit:** Create an app that allows users to input data, send it to the model endpoint, and display the prediction results.
 - **Azure Web App:** Alternatively, create a web app with a user-friendly form for making predictions.
 7. **Step 7: Monitoring and Management**
 - Set up logging and monitoring for your deployed model (e.g., using **Application Insights**).
 - Track prediction latency and other relevant performance metrics.
-

Deliverables

1. **Model Pipeline:**
 - Complete model pipeline from data preparation to deployment.
 - Experiment tracking with **MLflow** and results logging.
2. **Web Interface (Optional):**

- If using **Streamlit** or **Azure Web App**, provide a working interface that allows users to input data and receive predictions.
 - 3. **Report/Documentation:**
 - Provide a detailed report or presentation summarizing:
 - Dataset description and preprocessing steps.
 - Model selection and training process.
 - Hyperparameter optimization (if applicable).
 - Deployment steps and web interface implementation.
 - Results, model evaluation, and any improvements made.
 - 4. **Deployed Model:**
 - Provide access to the deployed model endpoint with example inputs and prediction results.
-

Additional Guidelines

- **Teamwork:** You can work individually or in teams of up to 3 members.
- **Datasets:** Choose a dataset from platforms like **Kaggle**, **UCI**, or others that fit your chosen problem type.
- **Optional Components:** Optuna and the front-end interface (Streamlit or Azure Web App) are optional but highly recommended for additional learning and project depth.